

The Feasibility of Stylometry as a Service

Josh Datko and Joseph Heenan

Drexel University

CS680: Privacy in the Electronic Society

Introduction

Maintaining anonymity online is increasingly difficult. While networks like Tor (?) will protect an Internet Protocol (IP) address at the network layer, it does not protect the content delivered over Tor. For example, suppose an author desires to provide controversial political commentary but wants to remain anonymous. In this case, she can protect her IP address with Tor and use end-to-end encryption to protect her message en-route, but she could be discovered through stylometry.

Stylometry, a form of authorship recognition, has been found to be extremely effective especially even at Internet-scale (?). Thus for those seeking to publish anonymous writings, not only can the delivery path reveal them, but their content can as well. Adversarial stylometry is the active process of manipulating writing to confuse the identity of the author. A set of tools, JStylo and Anonymouth (?), have been produced by the Privacy, Security, and Automation Lab (PSAL) at Drexel University to assist authors in anonymize their writing.

Currently the JStylo and Anonymouth integrated open-source project (JSAN) is available to download and can be installed on a machine that runs the Java Virtual Machine (JVM). While JSAN has been successfully demonstrated to be effective in assisting authors with adversarial stylometry (?), the constraint of downloading and installing the software on a machine powerful enough to run machine learning algorithms on data sets, limits their adoptability. We propose increase the availability and usability of JStylo and Anonymouth by offering them as a web service.

Migrating to a web service would provide JSAN with several benefits. This web service would provide adversarial stylometry capabilities to embedded devices like tablets or smart-phones which otherwise could not run the software natively. Since there would be no installation, besides a typical web-browser, a web-service would be better suited for time-critical situations where an author must quickly anonymize her document.

An example use-case is an author with a time sensitive document that she needs to release. Because of the sensitive nature, she wishes to apply adversarial stylometry to her

document. In this case, a computer with decent enough performance to natively run machine learning techniques may not be available. However, a web service is nearly ubiquitous.

Expanding Stylometric Research

A stylometric web service also expands research opportunities. This service, offered through a modern architecture, would allow researchers to easily perform experiments. An easily accessible tool allows for rapid verification of research results and greater community engagement.

In addition to convenient access, a web service can provide additional diversity in stylometric experimentation. For example, this web service supports multiple authors and provides a better tool for real-time, multiple author stylometry than a stand-alone application.

An effective technique in adversarial stylometry is imitation of another author's writing style. However, since many are not trained author imitators, this technique requires significant cognitive effort. Perhaps a more effective technique, would be to have multiple authors re-write the target work in their natural style. While there has been numerous studies into single author stylometric analysis, multiple authorship is more difficult. This web services provides a platform for multiple authorship research.

Challenges in a Web Service

A privacy enhancing technology that is designed to exist as a web service faces different challenges than an application. Attack vectors open up on the service from active denial-of-service attacks to passive sniffing. A privacy web service must take great care to ensure that service does not compromise its user's privacy. In Section ?? we provide our architecture and analysis the security and privacy.

1 Design

Anonymouth Architecture Review

The architecture of Anonymouth is that of a conventional desktop application. It uses a layered architecture with essentially two main components: the frontend Graphical User Interface (GUI) and the backend, consisting of the Weka machine learning toolkit and Jstylo. Jstylo is PSAL's stylometric analysis software. Since Anonymouth was designed as

a monolithic application, there are direct dependencies between the layers without the use of an adaptation layer.

This design is sufficient for Anonymouth in its current form. As a desktop application, the user can anonymize her and the document locally without any additional attack vectors.

Requirements for Web Service

The architecture for a web service imposes different requirements on an application. First and foremost, since Anonymouth is a privacy enhancing technology, the web service *must* not degrade the privacy protection of the application. This is a challenge when migrating from a desktop to web service since the desktop application does not use any networked communication. However, there is little benefit to using Anonymouth on the web if it compromises one's privacy.

Low-latency. The web service as implemented is near-real-time, i.e. low-latency, bounded primarily by the performance of Anonymouth. One key feature of the current implementation is the asynchronous processing of stylometric information, which allows a user to continually modify a text while anonymity information is calculated in the background. The latency of the response times for each stylometric analysis is on the order of 1 to 2 minutes in the default configuration (see Performance section for further details).

A high latency system would not be well suited for a web service. An example high latency system could simply entail sending an encrypted email to a third party who uses Anonymouth locally and returns the encrypted result. While this high-latency system seems trivial, note that a naive implementation will reduce the user's privacy. For example, email is very subject to traffic analysis (CITE) and encrypted email is easy to detect (CITE), therefore this act alone would raise suspicion. If this third party was offering a high-latency anonymization service for numerous people, the network graph would be easy to detect.

Architecture Qualities

There are several qualities of a system architecture that are considered foundational (CITE Lem Bass). These are: availability, interoperability, modifiability, performance, security, testability, and usability. In addition we add privacy as an architecture quality in our review.

Availability is the measure of the measure of the accessibility and reliability of the system, e.g. how well does it handle errors. Interoperability refers to the degree with which a system can interact with its components or other systems. Modifiability is important for large systems so that they can be easily maintained and scale. The response of the system to input is its performance. Security is the review of a system's ability to protect sensitive data and prevent unauthorized and undesired behaviour. Testability is an important factor not only in development but also in a deployed system to ensure that the system is working as designed. Usability is notoriously bad for security software (CITE How Johnny) and refers to the interaction design. Lastly, by privacy we mean how well a user's anonymity is protected.

In the following sections, we provide an analysis of each quality and define the requirements for a stylometric web service.

Availability. If the user can not access the web service, then the user is in the same situation as having a device that can not run Anonymouth locally. The entry point to the web service requires a web server. Realistically, this service will never see the traffic that Facebook or Twitter experiences, but the web server should be able to handle thousands of requests per day.

A common attack to block availability is a Denial-of-Service (DoS) attack. A DoS attack is akin to receiving a momentarily burst in traffic, i.e. the "Slashdot effect." The difference lies in the motivation for the sudden attention. Thus, a service that is resistant to DoS attacks also scales well to large second-derivative increases in traffic.

There are many implementation details to achieve availability that are beyond the scope of this paper. We attest that a stylometric service needs to be available and should consider the effects of sudden traffic increases.

Interoperability. The stylometric web service we implemented is interoperable with the modern web architecture. Specifically, it should be a service based on Representational State Transfer (REST)(?). REST web services define a framework that is client-server, has a uniform interface based on HTTP verbs, and supports a layered architecture. These features are in-line with the goals of a stylometric web service. Using a REST service, the client software can be a web browser, which is nearly ubiquitous on mobile devices.

Modifiability. One of the goals of the service is to provide a platform for stylometric research. In this sense it must be modifiable to accommodate different kinds of experiments including ones not currently anticipated. Software engineering best practices are applicable here mainly having well abstracted interfaces and modular components. As a RESTful web service, each endpoint should have a concise application programming interface (API) for a given feature set. For example, there should be a stylometry endpoint that can process text and return features. Other services, such as an adversarial stylometry service, should use the stylometric service via its REST API.

A web service has a distinct advantage here over the modifiability of the desktop software. One can deploy an updated version of the software by pushing the update to servers and instantly clients will have the new version. As a research platform, this is a very valuable feature. For example, split testing or A/B testing, is much easier to control and manage with a web service than packaged software.

Performance. The performance of the Stylometry Service as implemented is largely determined by that of the Anonymouth service that it uses to get stylometry suggestions. For an analysis of 1 500-word test document vs 23 documents from 4 'other' authors and using 8 training documents from the to-anonymize author, we recorded a total of 67 seconds required for the Anonymouth analysis, with less than 1 additional second used to update the client UI, on a 3.3 Ghz

Core i5 Ubuntu 12.04 Desktop machine, using 1 calculation thread for Anonymouth. These results suggest that any attempts at performance optimization should center around optimizing the performance of the Anonymouth configuration.

After each stylometric analysis is complete, the web service ensures all relevant anonymouth resources are closed. This allows the web service to be used over a long period of time (20 or more anonymization attempts have been tested) without memory leaks noted.

Further performance testing and optimization is required in order to support a true multi-user service. The authors would suggest using a web service based load test framework, such as Apache's JMeter or LoadUI, in order to fully test the service under various potential user loads.

Security. A number of attack vectors for the web application currently exist. We discuss those which seem most significant here below. However one key security implication of a web-delivered Stylometry Service is that it is expected to be delivered over the web, can be accessed in this configuration by anyone with knowledge of its unique URL (or .onion address). This allows adversaries greater opportunity to attempt to penetrate the security of the web application relative to a client application installed on a PC that is not internet connected.

It is expected that, before the service would be used in a real-world use case, a user would configure and deploy supporting databases associated with the login and SSL options available in the Play Framework.

Although web accessibility may be a significant risk from a usability standpoint, the anecdotal infrequency with which adversarial stylometry software is used even among privacy-conscious users as determined by conversations with colleagues of the authors suggests that delivery of this functionality via a web page (or even potentially via a browser plugin functional across multiple web pages) may be an appropriate trade-off to make.

Beyond the concerns mentioned above, a number of specific attack vectors remain for a stylometric web service in general and for our implementation in particular:

Man in the middle attacks: If a user accepts a rogue SSL certificate, it may be possible to create a "man in the middle" attack where the user believes he/she is communicating with a specific Stylometry Service, but is in fact communicating via an impersonated target.

Replay attacks: An eavesdropper may 'replay' a message to the webservice in an attempt to surreptitiously revert parts of the text to earlier versions. This may be possible as it is not documented whether the current implementation uses nonces when transmitting messages.

HTML5 Canvas attacks: The etherpad implementation uses the new 'canvas' element introduced in HTML5 extensively. Recently an attack has been documented (see Contextis Browser Timing Attacks white paper) whereby Firefox optimizations in the implementation of the requestAnimationFrame method have been used by to detect differences in link coloring across frames (allowing an adversary to determine if a particular stylometry service has been vis-

ited, for instance)

Note: We do not know if any of these vectors allows for a specific attack, but suggest they should be considered for any future implementation research in the area of stylometric web services using a similar architecture.

Testability. The web service UI is amenable to automated testing via web testing frameworks such as Selenium (although such tests were not written for this project). JUnit tests, in contrast, were written to test both service method paths implemented. The authors consider the fact that this application is delivered as a web application with a distinct service layer to be a positive from the standpoint of testability, as it allows users who may wish to modify the service the ability to expand an existing unit test suite.

Usability.

Privacy.

2 Security Analysis

3 Privacy and Anonymity Analysis

The privacy of a user accessing the web service is expected to be ensured via the use of TLS/SSL encryption for all data transmitted between the client and server. The anonymity of the user is expected to be promoted via the use of an Onion Routing protocol such as Tor, to ensure that the IP of the client cannot be directly linked to the IP of the service.

4 Implementation Details

We implemented a web application based on the Play Web Application Framework in Java. This web application had a reference to a modified version of JSAN which was forked from the PSAL Anonymouth repository (<https://github.com/psal/anonymouth>) on September 29, 2013. In addition the application has major dependencies on JQuery, Twitter Bootstrap (for page layout), and Etherpad (an open-source collaborative text editing web application and library). Etherpad (specifically Etherpad Lite is the version used here) in turn has major dependencies on node.js for building and running the backend service.

At a high level the implementation is as follows: A user browses to the Stylometry Service via a web browser. This causes the Play Framework web application to serve up a web page with an Etherpad control and instructions for the user to enter in the document they wish to anonymize. The etherpad control is implemented as an iframe which serves an Etherpad that is connected to the node.js server running the Etherpad backend (both servers should reside on the same domain, running on different ports). The page served by Play also contains a button which allows the author to specify when he or she would like to submit the document for anonymization to Anonymouth. Because the Anonymouth analysis process depends on a number of computationally-intensive algorithms (among them running SMO analysis in Weka as well as a Basic-9 feature stylometric analysis) we implemented this calculation asynchronously. When a user clicks the button to run an analysis, JQuery is used to issue a HTTP POST containing

in the message body a JSON-encoded representation of the current Etherpad text to an endpoint at `/getAnonymouthInfoJSON`. During the time that the Anonymouth analysis is underway, a spinner control from the `spin.js` library is used to indicate to the user that an analysis is underway. The fact that the request is made to Anonymouth asynchronously allows a user to continue editing the text while an analysis is running, allowing for a smooth user experience.

The web service as implemented currently supports only 1 author as the assumption is made that the author has pre-configured the "doc set" (also referred to as 'problem set' in Anonymouth source) that provides references as to the author's own writing style, as well as writing samples for other authors that will be used as a reference for writing styles that are different than that of the active author. The user can additionally configure the anonymouth properties file (`anonymouthprop.prop`) in the `jsan` resources directory to influence settings such as the feature set to use for stylometric identification as well as the number of threads to use and maximum number of features to analyze.

Once configured, the call to the `getAnonymouthInfoJSON` input will overwrite the `"test.txt"` file on disk that is configured in the provided problem set. This will ensure that changing the document in Etherpad forces an update to the test file and thus ensuring that the 'latest' document is always anonymized. Once the problem set has been updated, the 'Anonymouth library' is invoked. The 'Anonymouth library' is a forked version of the current Anonymouth application, which is designed to be used behind a Swing/AWT rich client GUI. A number of implementation challenges were present when trying to separate the core anonymization logic of Anonymouth from its GUI implementation - most specifically, although an 'engine' namespace exists which holds much of the anonymization logic, significant dependencies still existed between the 'engine' and the 'GUI' namespaces, specifically in the event handlers associated with the Editor screen and the `GUIMain` classes. Much of the work in this project consisted of reducing the coupling between these two namespaces.

As much of the configuration logic was handled at the GUI layer (via `ThePresident` and `GUIMain` classes) we took the approach of refactoring the GUI layer to work in a 'headless' mode, that is, without any windows launching. Long-term it would be preferable to have no dependencies on the GUI namespace at all for the Stylometry Service, however this could not be achieved in the available timeframe for this project. The web service thus invokes a modified version of the `ThePresident` class, which coordinates reading configuration files, loading the configured problem set, and using numerous libraries (JGAAP and Weka among them) to conduct an adversarial stylometric analysis. The results are returned by the web service in the form of a JSON-encoded list of words to remove, as well as an indication of 'changes needed' to the current text to reach an acceptable level of anonymity.

While this analysis has been running, the UI has been waiting on a response to the request made of the Stylometric Service. Once this request is returned, a request is made to `spin.js` library to stop the spinner, and JQuery is again used

to update a div on the page with the results of the adversarial stylometric analysis. The JQuery Promises interface was used to ensure that the UI correctly waited for these otherwise asynchronous requests to be completed in the right order.

The web application (bundled with all dependencies) can be downloaded by cloning <https://github.com/LRParser/anonymouth> and running the following command:

```
./bootstrap.sh
```

This command will copy the latest JSAN JAR files to the 'lib' directory of the Play Framework Web Application, start a Play Framework Server on port 9000, and start an Etherpad Lite server on port 9001 over `node.js`.

5 Comparison to Stand-alone Application

The stand-alone application has both advantages and disadvantages relative to the web service. To start with advantages: it contains more robust error handling, offers a GUI to create a problem set, and provides a richer editing experience for anonymized texts. Disadvantages concern ease of use (users must download and execute a JAR file) and supported platforms (it is unlikely that this JAR file will be supported on many mobile phones, for instance). The Stylometry Service, in contrast, is potentially more accessible than the standalone application: it can be accessed via mobile phones and tablets, for instance. Likewise application delivery may be simpler - nothing must be downloaded locally. Furthermore, the web application maintains a 'timeline' of changes made to a text and also allows for simple multi-party anonymization of documents.

Disadvantages to the web application include the fact that it offers fewer anonymization features (for instance, highlighting of specific text passages), instead focusing primarily on 'words to remove'. Furthermore, its security footprint is less well known, as it introduces a number of new dependencies / potential attack vectors such as JQuery, Twitter Bootstrap, Etherpad Lite, `node.js`, etc.

For a user likely to be using mobile devices or a casual user looking to quickly anonymize texts, we suggest that the web interface may be a better option. For a user facing a highly sophisticated adversary and needing detailed anonymization suggestions, we suggest that the stand-alone application may be a better option. We note that there is no technical reason that could keep the web application from offering all of the same anonymization suggestions as the standalone application, and thus it may be the case that in the future the web application could be recommendable to a wider range of users. standalone client

6 Single User Evaluation

Although time didn't permit detailed user testing of the web application, the authors did conduct one experiment with two users using Amazon Mechanical Turk. Users were given 45 minutes to anonymize a document. The title of the request was "Help re-phrase approx. 500 word texts as part of a research project". Specifically users were asked, 'Re-phrase the text in the below-linked document while keeping

the original meeting. Substitute synonyms, change punctuation (I am to I'm), split up / join sentences, etc. This is part of a research project to see how well different persons can make a text appear as if it wasn't written by the original author.'

The responses of users on Mechanical Turk (who were trying to anonymize a text without any supporting software) were compared with anonymization attempts by the authors of this paper against the same document. The anonymization attempts by the authors took place before reviewing the submissions from Mechanical Turk. For one sample paper, d 09.txt from the Brennan-Greenstadt corpus (Michael Brennan and Rachel Greenstadt. Practical Attacks Against Authorship Recognition Techniques in Proceedings of the Twenty-First Conference on Innovative Applications of Artificial Intelligence (IAAI), Pasadena, California, July 2009.), we noted that Anonymouth returned a "changes needed" score on the original document of 2.518, 2.677 for the users 'naive' attempt to anonymize over Mechanical Turk, and 2.467 for the attempt of one of the authors to anonymize using the web application created in this paper. Each of the two anonymization attempts for the 500-word document took approximately 15 minutes. Although these results are far from achieving statistical significance, they do suggest that it may be possible to use this software to achieve anonymization levels which surpass those of a 'naive' anonymization attempt.

7 Multiple Concurrent User Evaluation

TODO

8 Conclusion

There are two significant contributions as a result of this work. The first is a stylometric web service. This advances the usability and capabilities of JSAN and provide a framework for more expansive user testing in the future. The user testing results suggest that a web interface may offer usability advantages while still providing a meaningful increase in author anonymity. The second contribution is that this is the first framework for multi-author adversarial stylometry. Our results here suggest TODO.

Related Work

This proposal extends the work from Drexel's PSAL, who have made several contributions in this field. In (?), the authors present the Anonymouth framework for anonymizing writing. (?) confirmed adversarial stylometry is very effective at obfuscated authorship, but also presented techniques to hide the prescience adversarial stylometry. Lastly, (?) presented successful adversarial stylometry techniques through obfuscation, imitation, and translation and produced two corpora of texts.

Tor is one of the most successful anonymity services in operation (?). We intend to protect the anonymity of users by offering JSAN over a Tor hidden service however, there has been recent research de-anonymizing popular hidden services (?). We will provide a discussion on the design and

Milestone	Date
Document initial design	October 28
Finish porting JSAN to a web service	November 4
Add the Firebase editro	November 11
Add the third party feature	November 18
Conduct evaluation	November 25
Provide Paper and Presentation	December 2

Table 1: Project Milestones

experience of enabling a hidden service, but we will focus on this aspect.

Future Work

The scope of this project is designed to completed in one Drexel quarter, however we have several ideas for continued work. With a web service, we can offer real-time adversarial stylometry. Assume Alice, who has a blog, wants to anonymize her next blog post. She would connect to our web service, submit the URL for her blog and then write her next post with adversarial stylometry. This real-time data crawling at this scale is simply not feasible with stand alone software.

Now that a third party will be re-writing documents, we feel that it is important to obfuscate the content from the re-writer. One idea, is to mix the document, perhaps with other documents, prior to re-writing. The re-writer would be presented with an otherwise nonsensical document and would focus on re-write each sentence in her native style. The combined document, would contain n author styles, where n is the number of re-writing authors.

Lastly, for long term adversarial projects, we are curious on how often one has to re-train the model. We imagine that a web service could help monitor the change of a stylometric fingerprint over time.